

# COMS W4701: Artificial Intelligence

## Lecture 1b: Agents and Environments

Tony Dear, Ph.D.

Department of Computer Science

School of Engineering and Applied Sciences

# Today

---

- Task environments
- Agent programs
- Search problems

# Task Environments

---

- A rational agent is a *solution* to a **task environment** problem
- **PEAS**: Performance measure, environment, actuators, sensors
- Vacuum cleaner task environment
  - P: Cleanliness, power usage, time taken
  - E: The small grid world
  - A: Wheels to move, filter to clean
  - S: “GPS”, cleanliness sensor

# Example: Self-Driving Taxi

---

- **P:** Correct navigation; minimizing fuel, trip time, cost, traffic violations; maximizing comfort, safety, profits
- **E:** Roads, other traffic, pedestrians, customers, weather
- **A:** Acceleration, steering, braking, communication with others
- **S:** Cameras, GPS, speedometer, accelerometer, odometer, etc.

# Task Environment Properties

---

- Fully observable vs partially observable vs unobservable
  - Can agent sense all *relevant* information? Is internal state (memory) required?
- Single-agent vs multi-agent
  - Does maximization of performance measure depend on other agents' behaviors?
- Deterministic vs stochastic
  - Can we completely predict the next state of the environment?
- Episodic vs sequential
  - Do current decisions depend on past ones? Do current decisions affect future ones?
- Static vs dynamic
  - Does the environment change while the agent is thinking?
- Discrete vs continuous
  - Is number of states, actions, percepts, time, etc. finite?

# Examples of Environments

Environment	Partially / Fully Observable	Single- / Multi-Agent	Deterministic / Stochastic	Sequential / Episodic	Dynamic / Static	Continuous / Discrete
Vacuum cleaner world	Depends	Single	Deterministic	Sequential	Static	Discrete
Chess	Fully	Multi (adversarial)	Deterministic	Sequential	Static	Discrete
Self-driving car	Partially	Multi (cooperative)	Stochastic	Sequential	Dynamic	Continuous
Image classification	Fully	Single	Deterministic	Episodic	Static	Discrete


# Agent Design

---

- Understanding the task environment tells us how to design our agents
- The more difficult the task environment, the more complex the agent
- Partially observable env -> agent requires memory / state
- Multi-agent env -> agent requires tracking of other agents
- Stochastic env -> agent must consider multiple scenarios or contingencies
- Sequential env -> agent must consider past and future states
- Dynamic env -> agent must maintain model of the world

# Agent Functions

- An **agent function** maps percept sequences to actions
- Cleaner robot percepts: Current square; is dirty?
- Actions: Move left, move right, clean, do nothing




[A, IsClean]	Move right
[A, IsDirty]	Clean
[B, IsClean]	Move left
[B, IsDirty]	Clean
<hr/>	
[[A, IsDirty], [A, IsClean]]	Move right
[[B, IsDirty], [A, IsClean]]	Move right
[[B, IsClean], [A, IsClean]]	Do nothing
.....	.....



# Agent Programs

- An **agent program** is an *implementation* of an agent function
- Specifies *how* something is computed rather than *what* needs to be computed
- A given agent program is not a universal solution!
- Program usefulness depends on hardware, tractability
- Agent programs also depend on the desired *solutions*
- E.g., *optimal* wrt some utility, approximately optimal, *satisficing* (“good enough”), *probable*

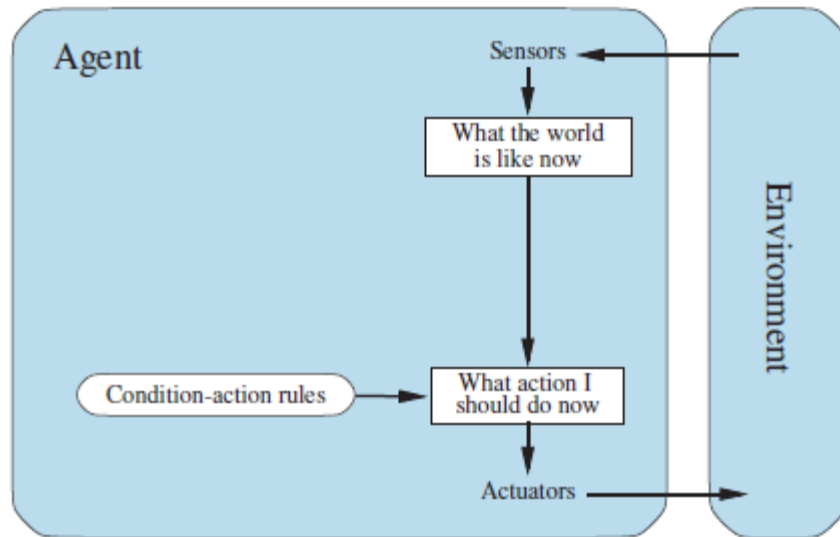


# Simple Reflex Agents

- Simple reflex agent: Use current percept only
- Percepts may be mapped to internal *states*
- Match state to condition-action (if-then) rules


```
function SIMPLE-REFLEX-AGENT(percept) returns an action
  persistent: rules, a set of condition–action rules

  state ← INTERPRET-INPUT(percept)
  rule ← RULE-MATCH(state, rules)
  action ← rule.ACTION
  return action
```



# Example: Vacuum Cleaner Robot

- Example: Vacuum cleaner robot as a reflex agent
- Use only the *current* (no past) percept
- What is its resultant behavior?




```
function Reflex-Vacuum-Agent([location, status]) returns action
  if status is dirty then return clean
  if location is A then return move right
  if location is B then return move left
  return do nothing
```

# Model-Based Reflex Agents

- What about partially observable environments?
- Maintain **internal state** and **transition model**
- May also have *sensor model* mapping percepts
- Use all information to *update* the state

```
function MODEL-BASED-REFLEX-AGENT(percept) returns an action
  persistent: state, the agent's current conception of the world state
               model, a description of how the next state depends on current state and action
               rules, a set of condition-action rules
               action, the most recent action, initially none

  state ← UPDATE-STATE(state, action, percept, model)
  rule ← RULE-MATCH(state, rules)
  action ← rule.ACTION
  return action
```



# Example: Vacuum Cleaner Robot

- We can provide our vacuum cleaner with *state* to *remember* its past

internalState = [*dirty*, *dirty*]

**function** Model-Based-Vacuum-Agent([location, status]) **returns** action


**if** status **is** dirty **then return** *clean*

**else** internalState[location] = *clean*

**if** internalState[B] **is** dirty **then return** *move right*

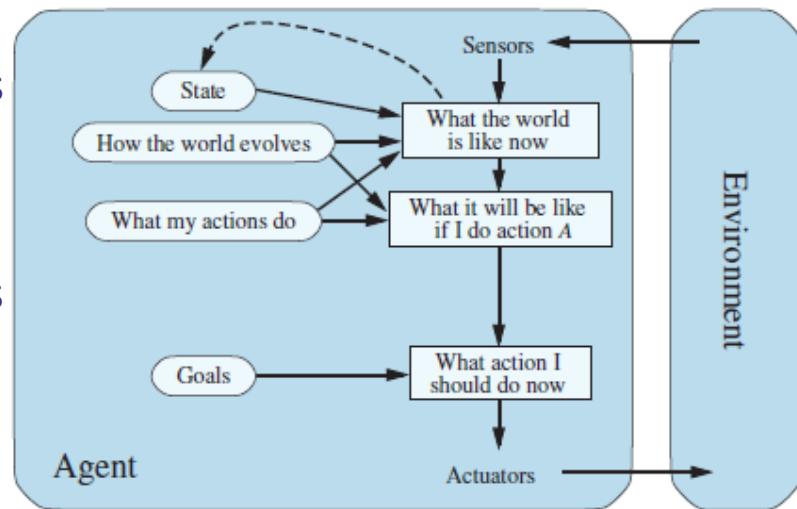
**if** internalState[A] **is** dirty **then return** *move left*

**return** *do nothing*



# Goal- and Utility-Based Agents

- Reflex agents are very rigid and predictable
- **Goal-based** agents try to achieve particular states
- Problems often solved using search and planning
- **Utility-based** agents can compare different states
- Utility functions map state to “desirability”
- Internalize the overall performance measure
- Utilities specify tradeoffs for competing goals
- Also useful in face of uncertainty



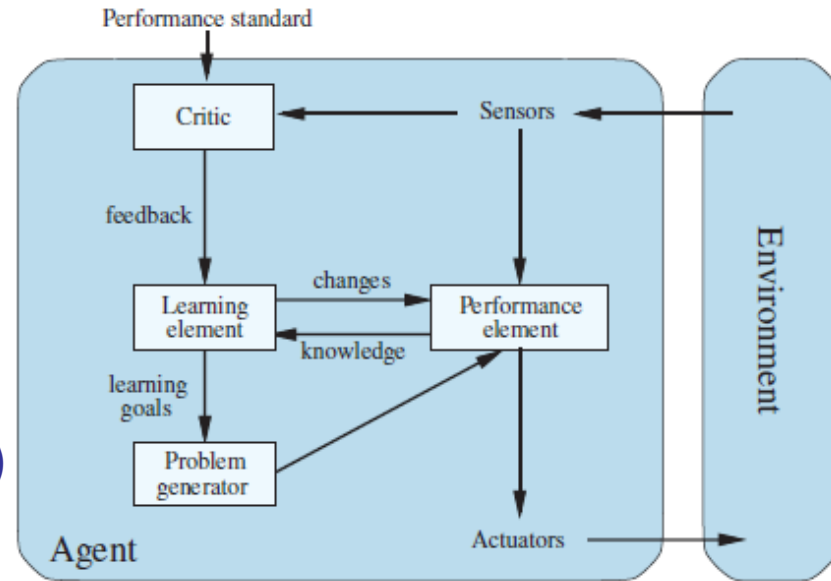
# Goals and Utilities

---

- Designing effective goal and/or utility functions is a hard problem
- May also be framed as objective, cost, or loss functions
- May be simple and concrete, or complex and abstract
- May be assigned explicitly, i.e., rewards in reinforcement learning
- May be implicit, i.e., perform similarly to training data examples
- Often captures tradeoffs between conflicting objectives

# Learning Agents

- **Learning agents** can be used to *create* or *improve* upon initial models in unknown envs
- A *learning element* retrieves knowledge from and then improves the *performance element*
- A *critic* evaluates the learning element according to a *performance standard* (measure)
- A *problem generator* suggests actions that can help gather new information and experiences





# Search Problems

---

- We start with task environments with the following properties:
- **Fully observable, single-agent, deterministic, static, discrete**
- We have “perfect” percepts and a known transition model
  
- *States* describe agent’s possible configurations as well as the goal
- Given problem description, the agent *searches* for a *state sequence* to the goal
  
- There may be *costs* incurred when going from one state to another
- An **optimal** solution minimizes total cost among all possible state sequences
- A **satisficing** solution may not be optimal, but is “good enough” with certain bounds


# Search Problems

---

- State space  $S$ : Set of descriptions of the agent and environment
- Initial state and one or more goal states
- Goal test  $S \rightarrow \{True, False\}$ , e.g.,  $isGoal(s_1) = False$
- Action set for each state, e.g.,  $Actions(s_1) = \{a_1, a_2, a_3\}$
- Transition model (function)  $S \times A \rightarrow S$ , e.g.,  $Result(s_1, a_1) = s_2$
- State-action cost function  $S \times A \rightarrow \mathbb{R}$ , e.g.,  $Cost(s_1, a_1) = 10$

# Example: Grid World Path Finding

- **State space:** Current coordinates of the rabbit
  - $S = \{(x, y) \mid x \in \{0,1,2,3\}, y \in \{0,1,2,3,4\}\}$
- **Goal test:**  $isGoal((3,3))?$
- **Actions:**  $Actions((x, y)) = \{Up, Down, Left, Right\}$
- **Costs:**  $Cost(s, a) = 1, \forall s, a$
- **Transition model:**  $Result((x, y), Up) = (x, y + 1), Result((x, y), Down) = \dots$ 
  - Can also account for walls and boundaries, e.g.  $Result((0,0), Left) = (0,0)$



# Example: $n$ -puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- **State:** Locations of all tiles and blank
- **Action:** 4 possible directions for the blank tile
- **Cost:** Each step taken costs 1
- **Goal test:** Is current state equal to goal state?

# More Search Problems

---

- Route-finding (e.g., vehicle navigation), robot navigation in the real world
- Touring problems (traveling salesperson)
- Layout and assembly sequencing problems
- Mathematical puzzles and proofs: Infinitely large state spaces!
- Knuth's conjecture (1964): Any positive integer can be obtained using a combination of factorial, floor, and sqrt operations on the input 4
- Ex:  $101 = \text{floor}\left((\sqrt{4!})!\right)$

# Summary

---

- PEAS descriptors define task environments and influence agent design
- Environment properties vary from easy to extremely challenging
- Agent programs may use current percept only, use a model, and/or try to achieve certain goals quantified by utilities
- Agent programs may also be created or improved via learning
- Problem-solving agents solve search problems to find discrete sequences of states and actions between initial and goal states