

# COMS W4701: Artificial Intelligence

## Lecture 2a: Uninformed Search

Tony Dear, Ph.D.

Department of Computer Science

School of Engineering and Applied Sciences

# Today

---

- State space graphs and search trees
- Uninformed search: DFS, BFS, UCS
- Variants: Iterative deepening, branch-and-bound

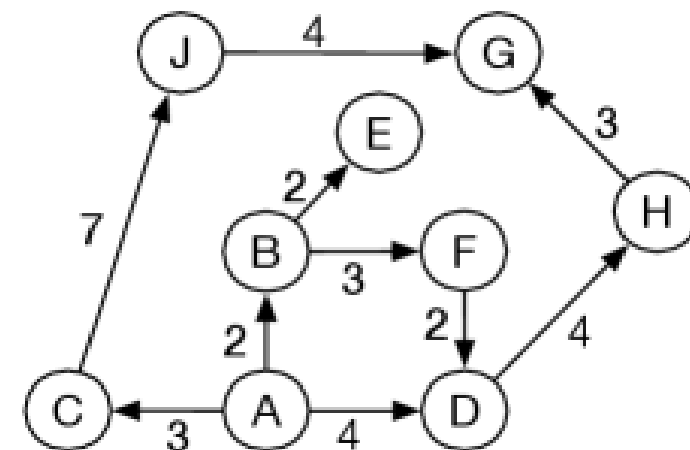
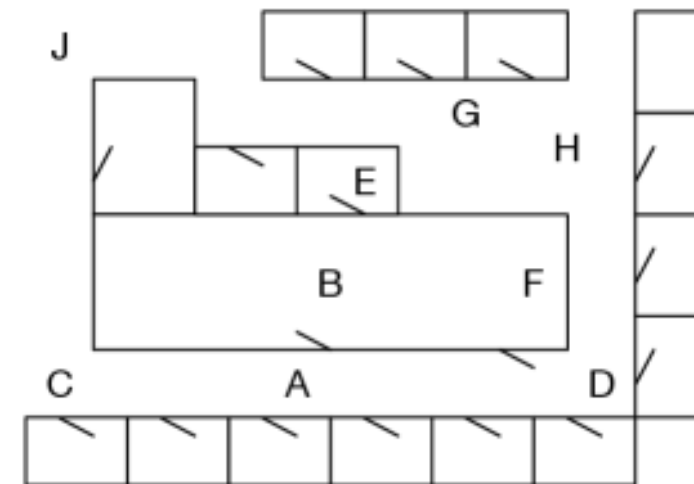
# Search Problems

---

- State space  $S$ : Set of descriptions of the agent and environment
- Initial state and one or more goal states
- Goal test  $S \rightarrow \{True, False\}$ , e.g.,  $isGoal(s_1) = False$
- Action set for each state, e.g.,  $Actions(s_1) = \{a_1, a_2, a_3\}$
- Transition model (function)  $S \times A \rightarrow S$ , e.g.,  $Result(s_1, a_1) = s_2$
- State-action cost function  $S \times A \rightarrow \mathbb{R}$ , e.g.,  $Cost(s_1, a_1) = 10$

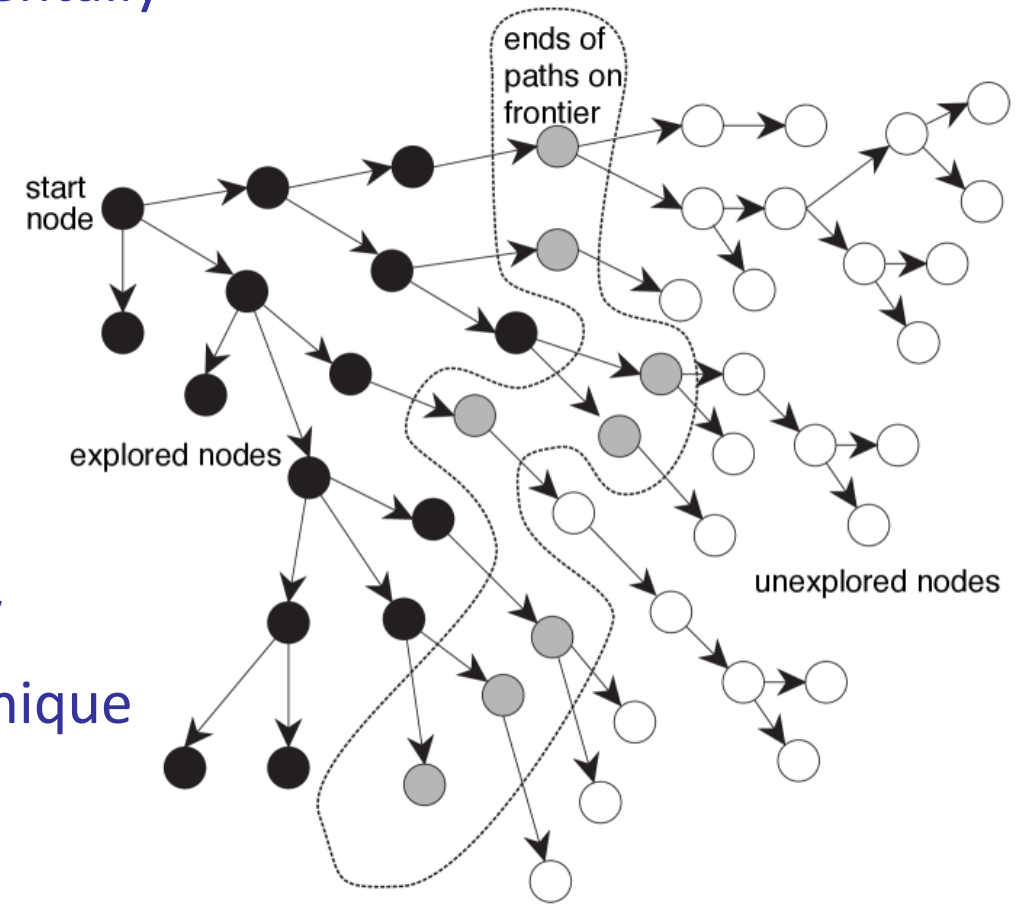
# State Space Graphs

- State space can be *abstracted* as a **directed graph**
  - States represented as vertices/nodes
  - Transitions represented as edges/arcs
  - Costs represented by edge weights
- 
- A solution is a *path* from initial to a goal state
  - State spaces can be very large or infinite, so we rarely build or store the full graph in memory



# Search Trees

- Most search algorithms explore paths by incrementally constructing a **search tree** from initial state
- Internal nodes have already been *explored*
- Leaf nodes lie in a **frontier** data structure
- Frontier nodes are candidates for *expansion*
- Important: Unexplored nodes are *not* in memory
- For explored nodes, tree structure records the unique path to each one from initial state



# Node Expansion

**function** EXPAND(*problem, node*) **yields** nodes

$s \leftarrow \text{node.STATE}$

**for each** *action* **in** *problem.ACTIONS(s)* **do**

$s' \leftarrow \text{problem.RESULT}(s, \text{action})$

$\text{cost} \leftarrow \text{node.PATH-COST} + \text{problem.ACTION-COST}(s, \text{action}, s')$

**yield** NODE(STATE= $s'$ , PARENT=*node*, ACTION=*action*, PATH-COST=*cost*)

- For each node that we explore in the search graph/tree, we *create* a new node for each *successor* and place them all in the frontier
- Can do so by looping over all available actions and looking at the results
- Also compute the *cumulative cost* of the new node—parent node cost + action cost
- Create successor node storing corresponding state, parent node, action, and cost

# Search Implementation

- We can now come up with a general *best-first-search* procedure
- Main idea: Repeatedly pop nodes out of and insert new nodes into frontier
- Run until we either find a goal or we run out of nodes to expand

**function** BEST-FIRST-SEARCH(*problem*, *f*) **returns** a solution node or *failure*

*node*  $\leftarrow$  NODE(STATE=*problem*.INITIAL)

*frontier*  $\leftarrow$  a priority queue ordered by *f*, with *node* as an element

Initialize current node to contain the initial state, frontier as a queue data structure

**while not** IS-EMPTY(*frontier*) **do**

*node*  $\leftarrow$  POP(*frontier*)

**if** *problem*.IS-GOAL(*node*.STATE) **then return** *node*

“Late” goal test (upon popping from frontier)

**for each** *child* **in** EXPAND(*problem*, *node*) **do**

Node expansion

            add *child* to *frontier*

**return** *failure*

# Frontier and Reached States

---

- The *search strategy* determines the ordering for popping frontier nodes
- Different strategies yield different efficiencies and guarantees
- Frontier implementation examples: LIFO (stack), FIFO (queue)
- Priority queue using general **evaluation function**  $f(n)$
- We may also want to *prune* the search space to make it smaller
- There may be multiple paths to a state; should check whether a state has already been *reached* before placing into frontier



# Search Implementation with Reached

**function** BEST-FIRST-SEARCH(*problem*, *f*) **returns** a solution node or *failure*

*node*  $\leftarrow$  NODE(STATE=*problem*.INITIAL)

*frontier*  $\leftarrow$  a priority queue ordered by *f*, with *node* as an element

*reached*  $\leftarrow$  a lookup table, with one entry with key *problem*.INITIAL and value *node*

**while not** IS-EMPTY(*frontier*) **do**

*node*  $\leftarrow$  POP(*frontier*)

**if** *problem*.IS-GOAL(*node*.STATE) **then return** *node*

**for each** *child* **in** EXPAND(*problem*, *node*) **do**

*s*  $\leftarrow$  *child*.STATE

**if** *s* is not in *reached* **or** *child*.PATH-COST < *reached*[*s*].PATH-COST **then**

*reached*[*s*]  $\leftarrow$  *child*

add *child* to *frontier*

**return** *failure*

*reached* may be implemented as a lookup table (dictionary) mapping states to nodes

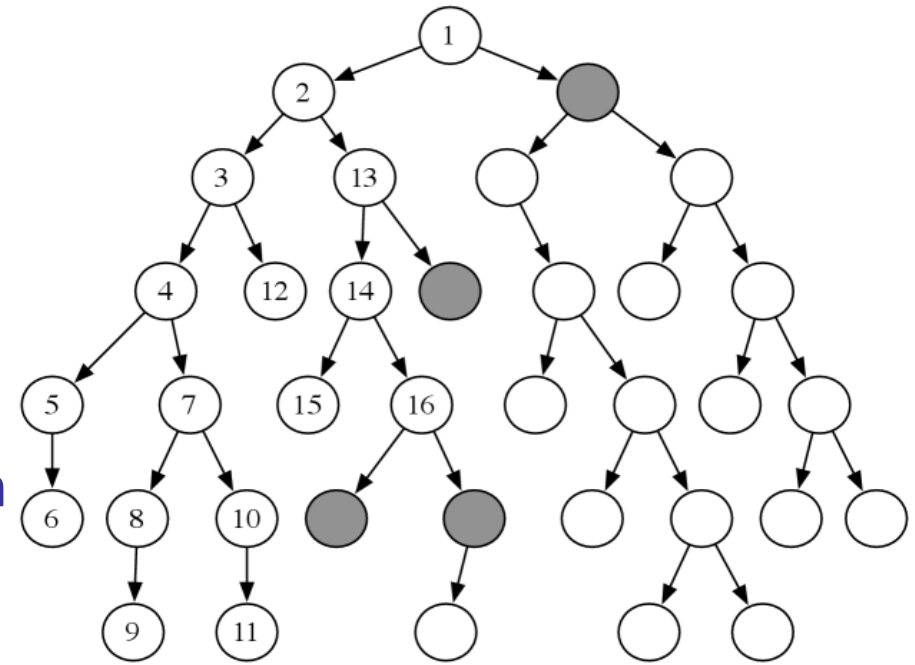
Two scenarios for adding node to frontier:

1. Previously unencountered (not in *reached*)
2. Already encountered, BUT found cheaper path

When adding a node to frontier, also add its state to *reached*

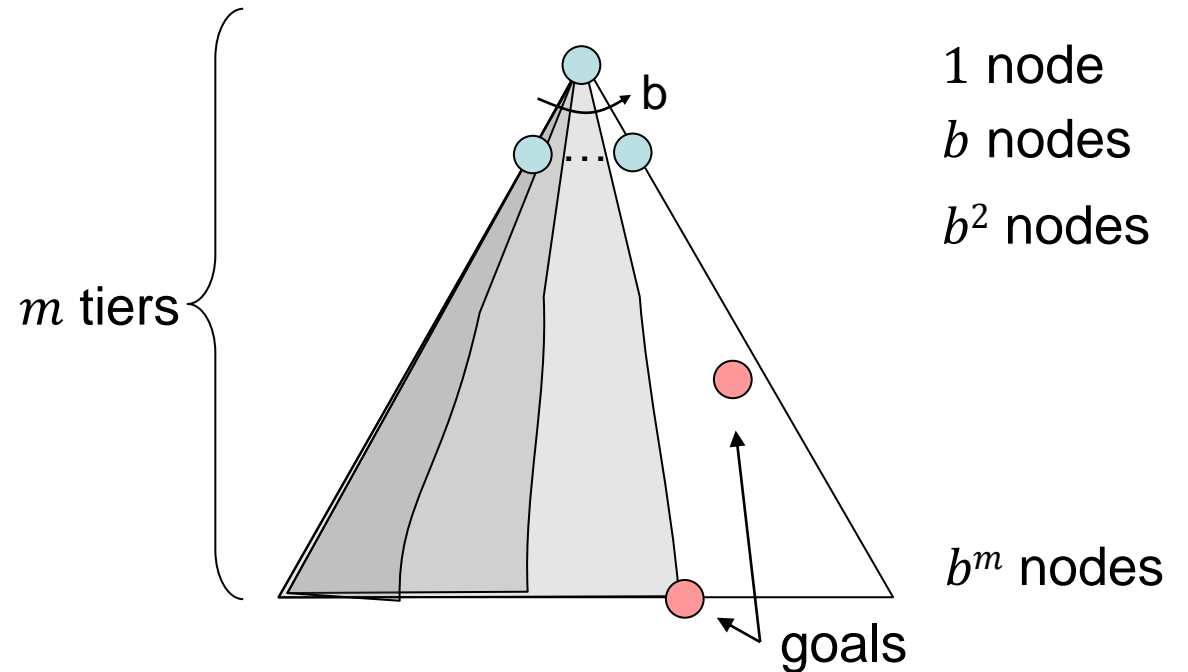
# Depth-First Search

- Idea: Expand the *deepest* node in the frontier, disregarding action costs
- Implement frontier as a stack (LIFO) or priority queue where  $f(n) = -depth(n)$
- Behavior ends up searching each path to completion
- If no goal, *backtrack* and proceed along a new path
- DFS is not **optimal**: no guarantee of least-cost solution
- DFS is not **complete**: no guarantee of finding a solution if one exists or correctly reporting failure otherwise



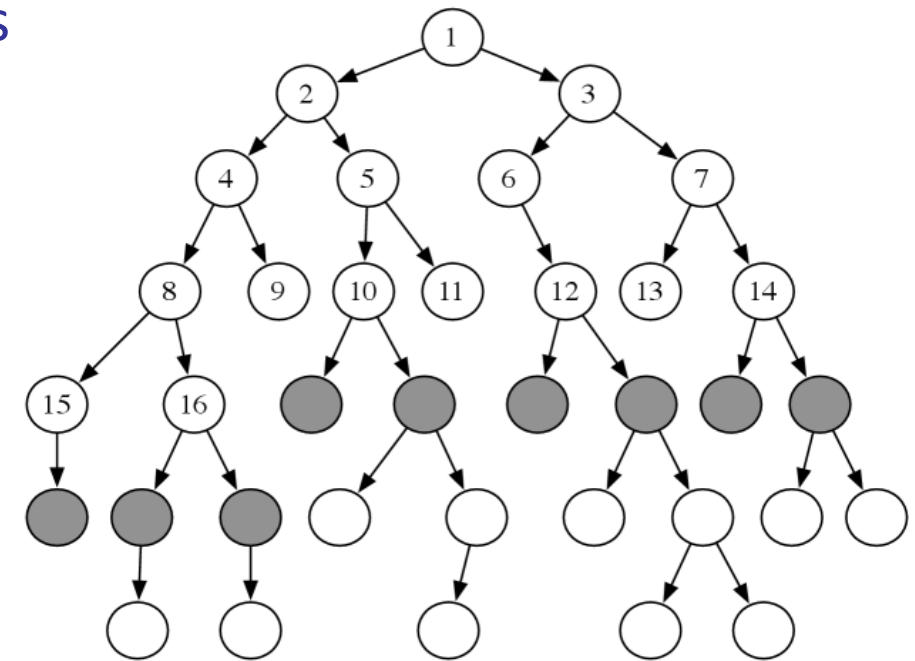
# DFS Complexity

- Suppose search tree has *branching factor*  $b$  and *max depth*  $m$
- **Time complexity:** (Worst case) number of expanded nodes  $O(b^m)$
- **Space complexity:** (Worst case) number of frontier nodes in memory  $O(bm)$ 
  - One node in each of  $m$  tiers with  $b$  successors each
- Good use cases for DFS when space is restricted, or many goals exist and we just want to find one fast



# Breadth-First Search

- Idea: Expand the *shallowest* node in the frontier, disregarding action costs
- Implement frontier as a FIFO queue or priority queue where  $f(n) = \text{depth}(n)$
- BFS is **complete**; will eventually find a goal if one exists
- BFS returns shallowest solution (e.g., at depth  $d$ )
- **Optimal** if action costs are uniform
- Worst-case time and space complexities are  $O(b^d)$
- Good use cases for BFS when problem / state space are small, or solution is close to start



# Comparison: DFS and BFS

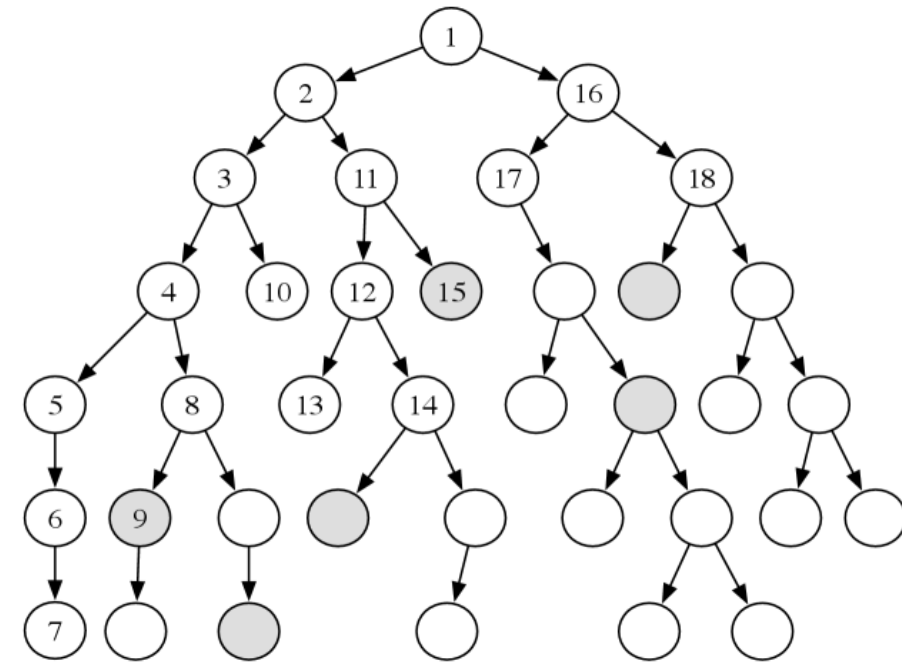
- DFS and BFS are both implementations of best-first search
- Frontier nodes are popped out *deepest* and *shallowest* first, respectively
- Both have exponential time complexity; BFS has exponential space complexity while DFS space is linear in search tree depth
- Neither is optimal in general; BFS only guaranteed to return shallowest solution
- Small optimization: “Early” goal test
- Check upon *insertion* into frontier

```
while not IS-EMPTY(frontier) do  
    node ← POP(frontier)  
if problem.IS-GOAL(node.STATE) then return node  
    for each child in EXPAND(problem, node) do  
        if problem.IS-GOAL(child.STATE) then return child
```



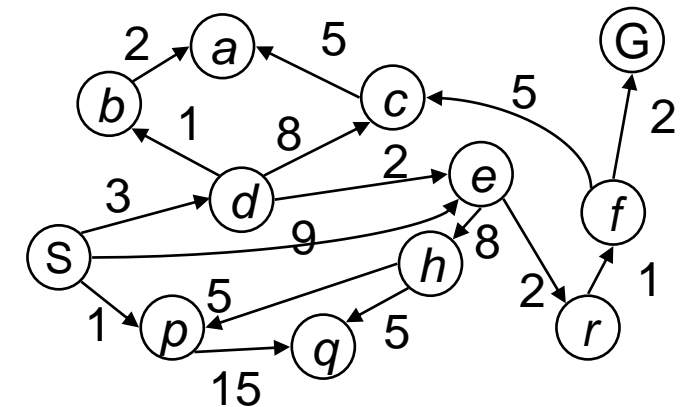
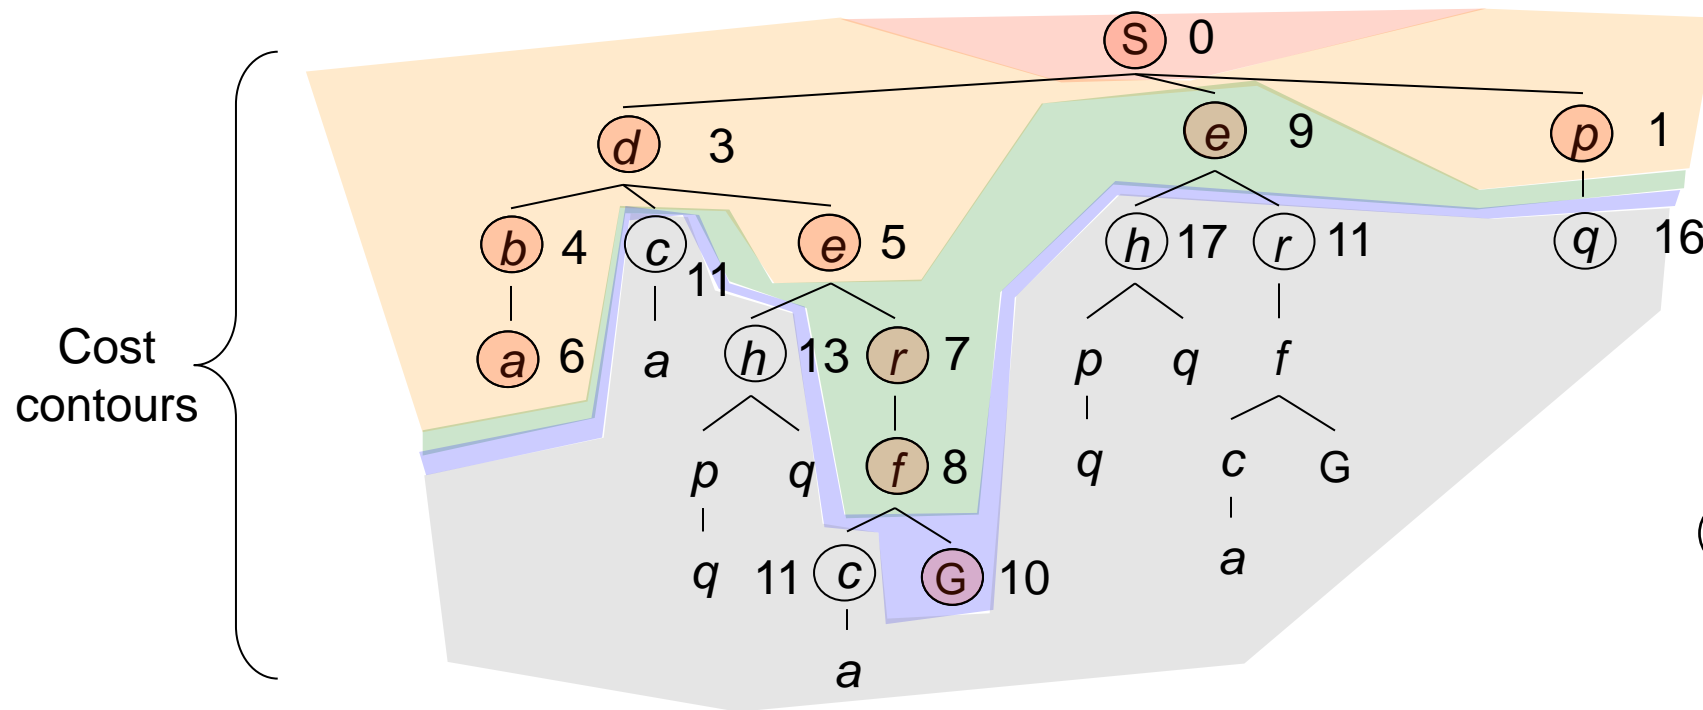
# Branch-and-Bound

- We can also set an *adaptive* depth limit based on best solution found so far
- Best use case when there are multiple solutions and we want to find all of them
- First run DFS as usual until a goal  $G$  is found
- Save the solution and continue DFS, but now with depth or cost limit equal to that of  $G$
- If we find better solution, save it and update the bound
- Search until frontier is empty and return best solution
- Can also be combined with iterative deepening!



# Uniform-Cost Search (Dijkstra)

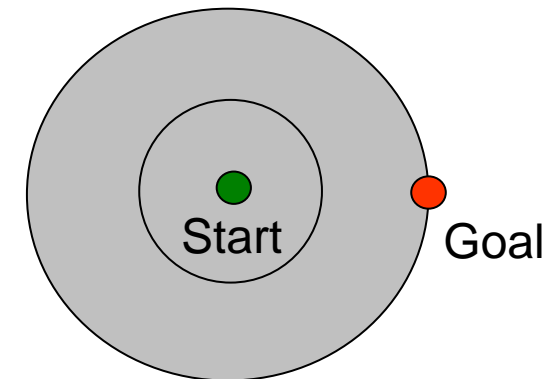
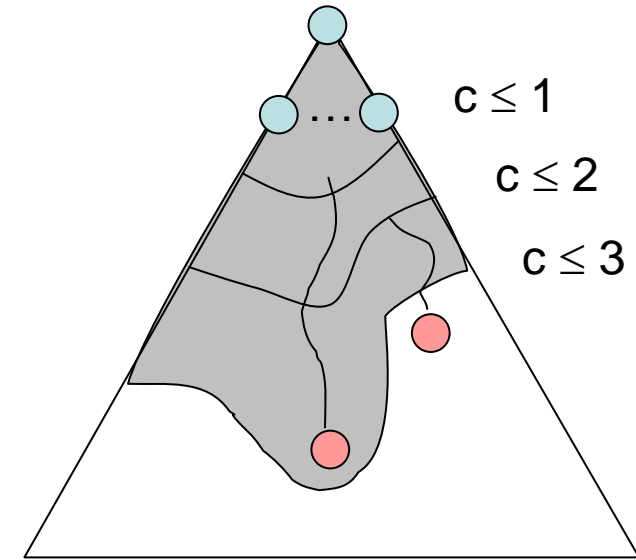
- To find an optimal solution, we expand nodes in order of *cumulative path cost*
- Implement frontier as priority queue with evaluation function  $f(n) = cost(n)$
- First goal that is found is guaranteed to be the one with lowest cost





# UCS Properties

- Let  $C^*$  be the cost of optimal solution
- Let  $\epsilon$  be lower bound on all possible costs
- $1 + \lfloor C^* / \epsilon \rfloor$  is the max depth to traverse before finding optimal solution
- **Time and space complexity:**  $O(b^{1+\lfloor C^* / \epsilon \rfloor})$
- UCS is both **complete** and **optimal**



# Summary

---

- Problem-solving agents solve search problems to find discrete sequences of states and actions between initial and goal states
- Search problems represented as graphs and/or trees, which we try to represent and search in a systematic and efficient manner
- Various search strategies, each with different tradeoffs in time complexity, space complexity, completeness, optimality
- Algorithms: DFS, BFS, iterative deepening, branch-and-bound, UCS