

COMS W4701: Artificial Intelligence

Lecture 2c: Local Search

Tony Dear, Ph.D.

Department of Computer Science
School of Engineering and Applied Sciences

Today


- Local search in discrete spaces
- Hill-climbing
- Simulated annealing

Local Search Algorithms

- Search algorithms covered so far may traverse all possible states
- If space is finite and solution exists, (most) algorithms guaranteed to find it
- Systematic search do not work well in large, infinite, or continuous spaces
- In some problems, we care more for final state than the actual trajectory
- **Local search** approaches only keep track of current state
- Use very little (constant) memory and can often find solutions in large spaces in practice, though with fewer guarantees


Example: n -Queens

- CSP solvers incrementally generate consistent assignments until complete
- Now consider a state space of all complete assignments of n queens, both consistent and inconsistent
- **Random sampling:** Randomly generate states until a valid one is found
- **Random walk:** From the current state, randomly generate a *successor state* by modifying a single queen assignment
- Repeat until valid state is found



Example: n -Queens

- Both random sampling and random walks are *complete* with infinite time
- No upper limit on time needed, may be very slow in practice
- **Iterative best improvement** selects successor states by minimizing (greedy descent) or maximizing (hill climbing, greedy ascent) an *evaluation/objective function*
- E.g., number of *conflicts* h in current state with valid solution $h = 0$




Iterative Best Improvement

```
function HILL-CLIMBING(problem) returns a state that is a local maximum  
  current  $\leftarrow$  problem.INITIAL  
  while true do  
    neighbor  $\leftarrow$  a highest-valued successor state of current  
    if VALUE(neighbor)  $\leq$  VALUE(current) then return current  
    current  $\leftarrow$  neighbor
```

- Hill climbing and other iterative methods may get stuck at **local optima**, states at which which no neighbor is better than the current one
- We prefer **global optima**, but no guarantee that we can find one
- Hill climbing is neither optimal nor complete


State Space Landscapes

- Can think about search space as a “landscape” of different features
- In addition to local/global optima, there may be shoulders or “flat” optima
- We may include *sideways* moves to equally valued successors, but need to avoid getting stuck in endless cycles
- We may include *random* moves to escape local optima



Example: n -Queens

- Current board has $h = 1$ conflict
- All neighboring states are the same or worse
- To escape this flat local minimum, we can allow sideways moves to other states with $h = 1$
- But limit number of sideways moves to avoid cycles
- For 8-queens, success prob without sideways moves is ~14%, with up to 100 sideways moves ~94%



Stochastic Local Search

- We can add randomness to iterative search to expand exploration
- E.g., take a *random step* to a worse successor
- Or do a *random restart* and move to a completely different state
- *Stochastic* hill-climbing: Instead of selecting best successor, select one based on a probability distribution assigned to all successors
- *First-choice* hill-climbing: Instead of evaluating all possible successors, just select a random (e.g., the first) one that improves upon current state

Simulated Annealing

- **Simulated annealing:** Generate random successor and move to it if better
- Otherwise, we *may* move to a worse successor with some probability
- Probability inversely proportional to how “bad” the successor is as well as time
- Time to probability mapping is determined by the *temperature* T

function SIMULATED-ANNEALING(*problem*, *schedule*) **returns** a solution state

current \leftarrow *problem*.INITIAL

for $t = 1$ **to** ∞ **do**

T \leftarrow *schedule*(*t*)

if $T = 0$ **then return** *current*

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow \text{VALUE}(\text{current}) - \text{VALUE}(\text{next})$

if $\Delta E < 0$ **then** *current* \leftarrow *next*

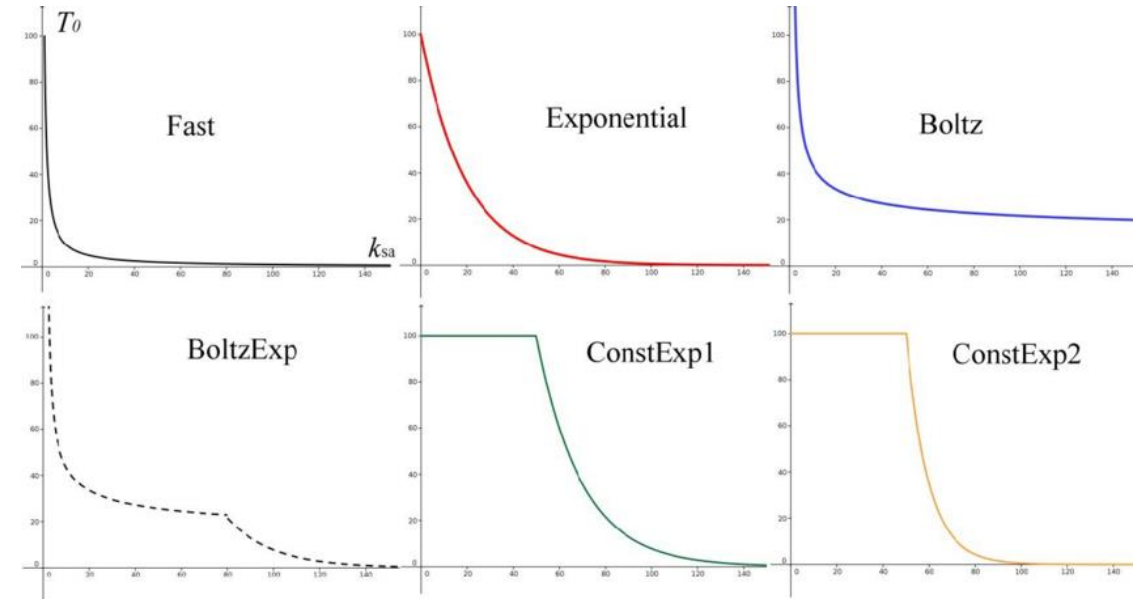
else *current* \leftarrow *next* only with probability $e^{-\Delta E/T}$

T typically nonnegative,
decreases over time

Move to *next* if it is better than
current, otherwise do so with
probability that decreases the
worse the *next* state is

Simulated Annealing

- $current \leftarrow next$ only with probability $e^{-\Delta E/T}$
- The worse the successor, the larger the ΔE , and the lower the probability of moving
- T typically follows a *schedule* function
- Schedule choice is problem dependent
- High or constant value when starting, higher probability of exploring worse states
- Temperature decreases to 0 over time and algorithm will only improve current state



Local Beam Search

- Instead of starting with and maintaining one state, start with k states
- Each one generates a successor, and we keep the k best ones
- Unlike (regular) local search with k random restarts, local beam search quickly shares useful information among the parallel search threads
- Disadvantage: Potential lack of diversity in states in the threads
- *Stochastic* beam search: Randomly choose k successors to keep in frontier, with likelihoods proportional to state values

Summary

- Local search algorithms: Good for problems that only require goal states, not entire paths; very memory-efficient
- No systematic search of state space; fewer optimality and completeness guarantees; tuning of parameters often required
- Hill-climbing: Move toward neighboring states that look better
- Simulated annealing: Occasionally allow moves toward worse states