

COMS W4701: Artificial Intelligence

Lecture 9b: Probabilistic Learning

Tony Dear, Ph.D.

Department of Computer Science

School of Engineering and Applied Sciences

Today

- Machine learning foundations
- Bayes net learning
- Laplace smoothing
- Expectation-maximization

Learning

- Given data generated by an unknown function f , a learning agent seeks to find a **hypothesis** \hat{f} to “best” approximate f
- f may be a function of input variables / predictors / features
- Alternatively, f may be a joint probability distribution
- \hat{f} can be used for *prediction* (e.g., classification, regression) of data inputs
- \hat{f} can be used for *inference* of relationships between inputs and outputs
- \hat{f} can be used to find *patterns* in input data (e.g., clustering)

Methods for Learning

- **Parametric methods** first specify a class of functions (e.g., linear combination of features) and then estimate their *parameters*
- In Bayes nets, maximum likelihood can be used for **density estimation**
- **Nonparametric methods** do not make assumptions about the model class, can potentially fit many more functions (e.g., k-nearest neighbors)
- May be more flexible, but may be harder to learn or less interpretable
- Learning may be supervised, unsupervised, or somewhere in between depending on whether output values are available

Performance Evaluation

- We first specify hypothesis space in which we search/optimize
- We then train a model on using a given set of *training data*
- Next, find ways to slightly modify model to generalize it
- Methods: Cross-validation, hand-tuning, grid search, etc.
- Use *validation data* to tune any **hyperparameters**
- Tuning lowers training accuracy, increases validation accuracy
- Final model evaluation is done on a new set of *test data*

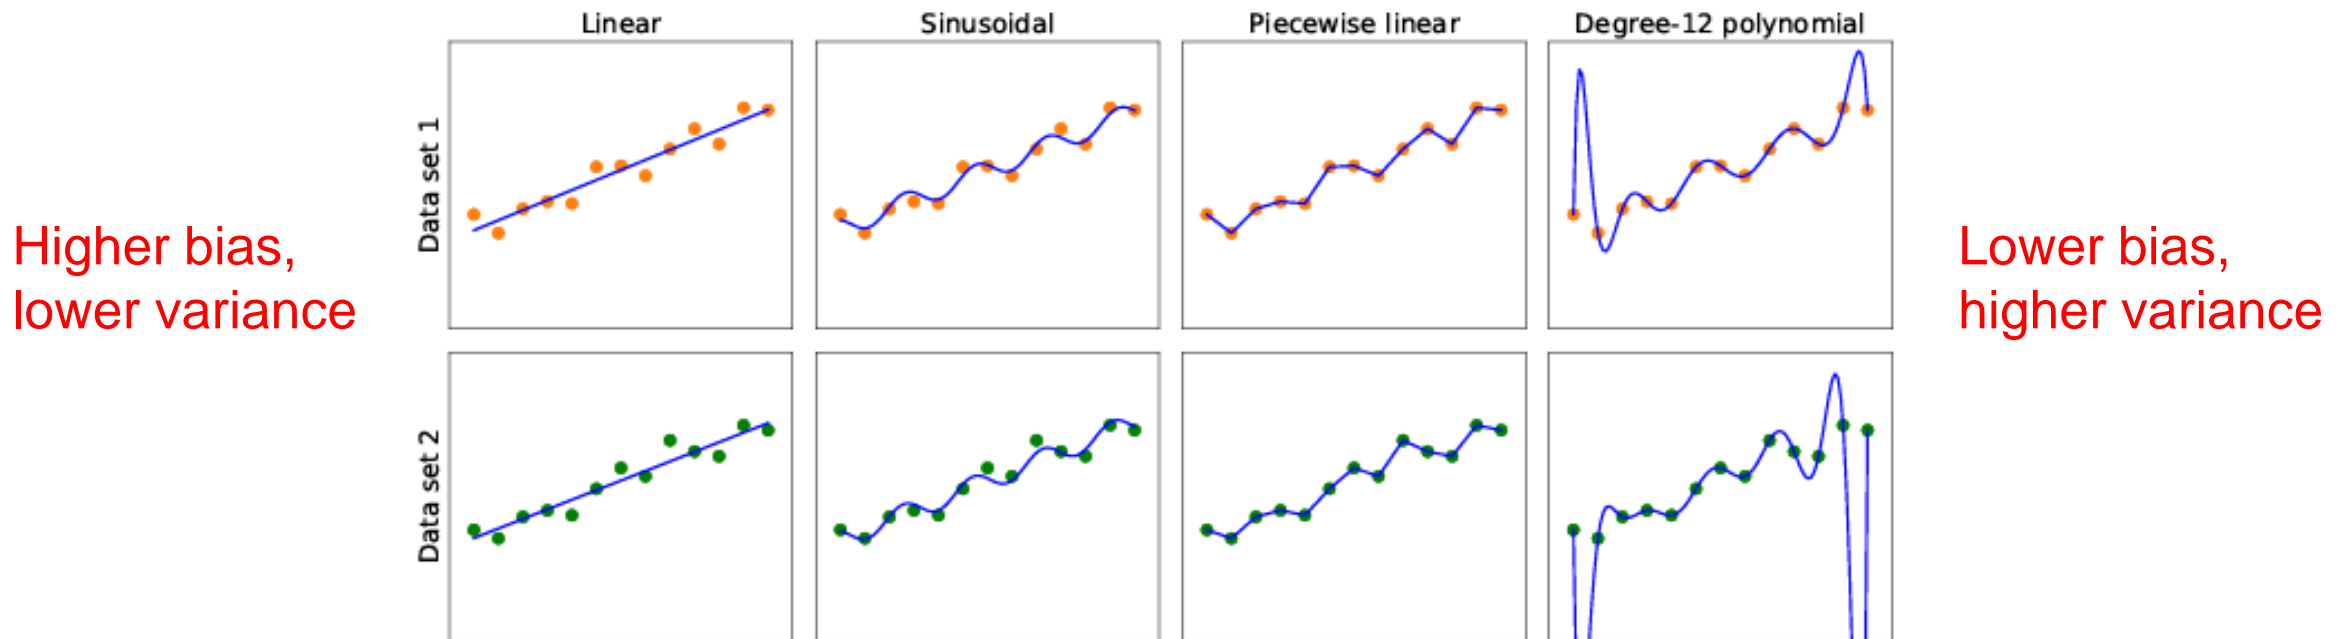
Training
Data

Validation
Data

Test
Data

Bias-Variance Tradeoff

- More complex models will generally better fit the given data, resulting in higher accuracies, lower error, less *underfitting*, and lower **bias**
- Tradeoff: Models that are too flexible may *overfit* to noisy or unrepresentative training data, resulting in greater **variance** in the learned models given different data sets



Other Learning Considerations

- Models with higher complexity can be used for greater variety of problems
- E.g., deep neural networks are very successful at hard vision/NLP tasks
- Tradeoff: More data required, more computation for learning a good model
- Higher complexity also generally leads to lower interpretability
- May be good for prediction, but not so good for inference and explainability
- When data is high-dimensional, data sets become very sparse, leading to high variance
- Feature selection and dimensionality reduction techniques can help
- Learning methods must also be able to handle imperfect, noisy, and missing data

Bayes Net Learning

- Suppose we want to learn the *parameters* θ (probabilities) of a Bayes net
- Hypothesis space is the set of all possible CPTs / joint distributions
- **Maximum-likelihood** of independent and identically distributed data $\mathbf{d} = (d_1, \dots, d_N)$:

$$\max_{\theta} \Pr(\mathbf{d}|\theta) = \max \prod_{i=1}^N P(d_i|\theta)$$

- $P(d_i|\theta)$ is the Bayes net joint probability of the data point d_i
- We can perform optimization to solve this, but solution is (again) intuitive!
- MLE parameters are just *frequencies*, and each CPT can be learned *independently*

Naïve Bayes Models

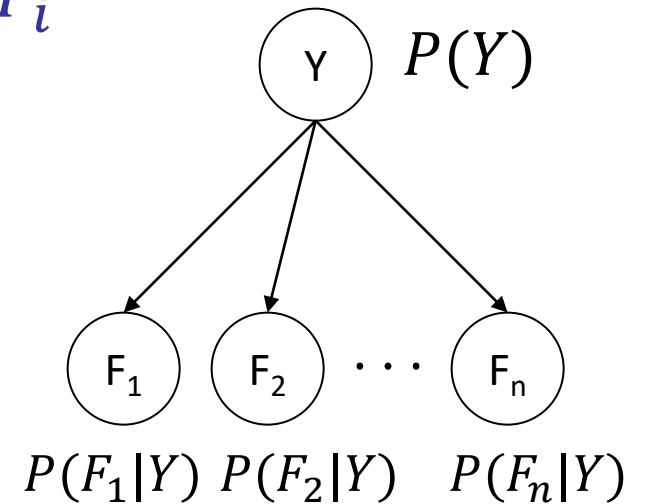
- **Naïve Bayes** model: Fork structure with single class or label variable Y
- Each label is correlated with many different features F_i

- Joint probabilities:
$$P(y, f_1, \dots, f_n) = P(y) \prod_{i=1}^n P(f_i|y)$$

- A common inference task is **classification**

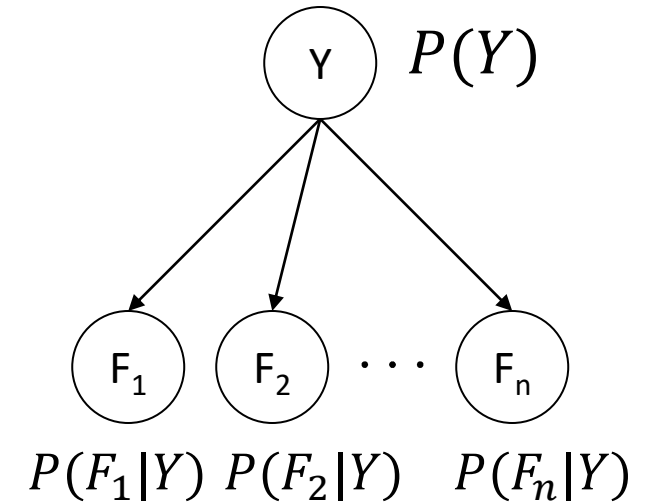
- Suppose are given an input *feature vector* (all f_i)

- Most likely class label:
$$y = \operatorname{argmax}_y P(y|f_1, \dots, f_n) = \operatorname{argmax}_y P(y) \prod_{i=1}^n P(f_i|y)$$



Example: Spam Filtering

- $P(Y)$ = Spam prior, or likelihood that generic email is spam
- $P(F_i|Y)$ = Conditional probabilities that a word F_i appears or doesn't appear in a spam or ham email
- **“Bag-of-words”** model: Each word is treated independently, no dependency on other or nearby words



	$P(Y)$	Word	Hey	would	you	like	to	lose	weight
Spam	1/3	$P(f \text{spam})$	0.00002	0.00069	0.00881	0.00086	0.01517	0.00008	0.00016
Ham	2/3	$P(f \text{ham})$	0.00021	0.00084	0.00304	0.00083	0.01339	0.00002	0.00002

- **Prediction:** $\operatorname{argmax}_{\text{spam,ham}} \left\{ P(\text{spam}) \prod_{i=1}^n P(f_i|\text{spam}), P(\text{ham}) \prod_{i=1}^n P(f_i|\text{ham}) \right\} = \text{spam}$

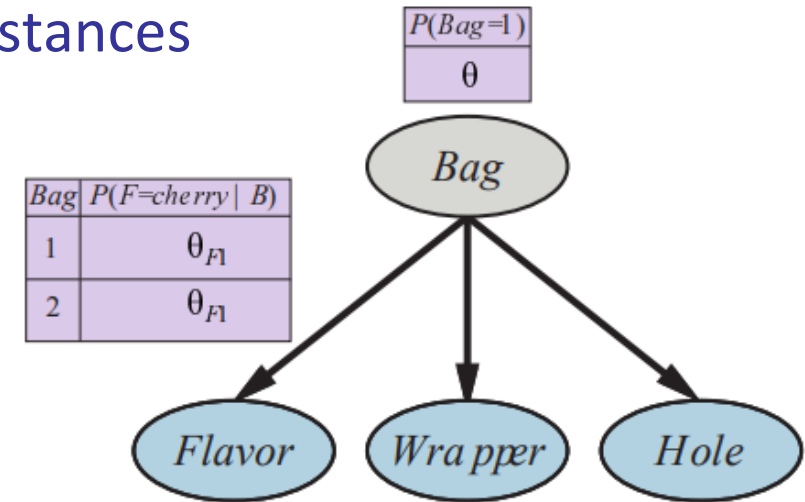
Learning Naïve Bayes Models

- Example: Candy bag prediction given candy features
- Suppose data \mathbf{d} has b_1 bag 1 instances and $n - b_1$ bag 2 instances

- Data likelihood for *Bag* class: $P(\mathbf{d}|\theta) = \theta^{b_1}(1 - \theta)^{n-b_1}$
- It is often convenient to work with *log likelihood function*:

$$\begin{aligned} L(\mathbf{d}|\theta) &= \log P(\mathbf{d}|\theta) = \log \theta^{b_1} + \log(1 - \theta)^{n-b_1} \\ &= b_1 \log \theta + (n - b_1) \log(1 - \theta) \end{aligned}$$

- To maximize wrt θ , solve $\frac{dL}{d\theta} = 0$: $\frac{dL}{d\theta} = \frac{b_1}{\theta} - \frac{n - b_1}{1 - \theta} = 0 \quad \Rightarrow \quad \theta = \frac{b_1}{n}$



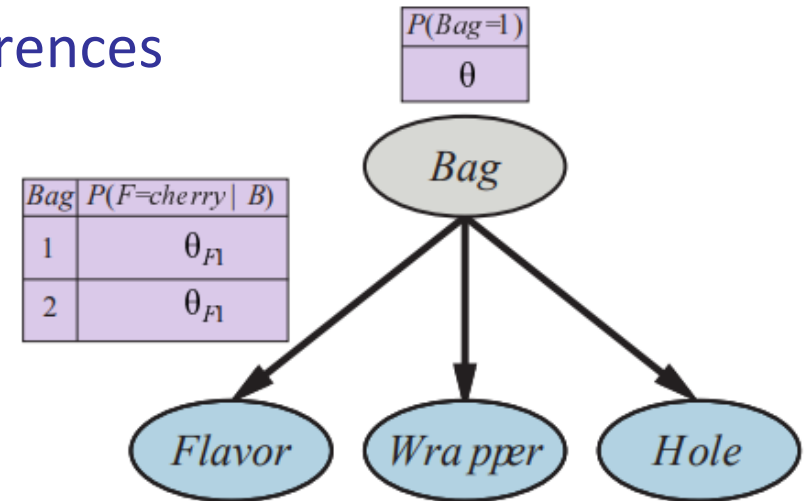
Learning Naïve Bayes Models

- Can write likelihoods for each feature and optimize similarly
- Results always come down to simply counting data occurrences

- Maximum-likelihood parameter estimates:

$$\theta = \frac{N(bag = 1)}{N(samples)} \quad \theta_{F1} = \frac{N(bag = 1, flavor = cherry)}{N(bag = 1)}$$
$$1 - \theta = \frac{N(bag = 2)}{N(samples)} \quad \theta_{F2} = \frac{N(bag = 2, flavor = cherry)}{N(bag = 2)}$$



- Suppose training data contains no instance of $(Bag = 1, F = cherry)$ candies
- When predicting on future instances, the classifier will give $P(Bag = 1) = 0$ for any input containing $F = cherry$, regardless of other features!



Laplace Smoothing

- **Data fragmentation** in Bayes nets: Each parameter is estimated using only a subset of the training data, leading to *overfitted* estimates
- Problem worsens as the CPTs become larger (higher dimensionality)
- **Laplace smoothing:** Reduce overfitting by adding a “virtual count” α

$$\widehat{P(X)} = \frac{n(x) + \alpha}{N + \alpha|X|}$$

$\alpha \rightarrow 0$ 	$\widehat{P(X)} = \frac{n(x)}{N}$	Original maximum likelihood estimator (average)
$\alpha \rightarrow \infty$ 	$\widehat{P(X)} = \frac{1}{ X }$	Uniform prior over all possibilities of X (ignore all samples)

Example: Laplace Smoothing

- Suppose we have the following training data set:
- Estimates of class priors: $P(+y) = P(-y) = 0.5$
- MLE would give $P(+f_3|+y) = 0$ and $P(-f_2|-y) = 0$
- Suppose we use $\alpha = 1$ smoothing instead:

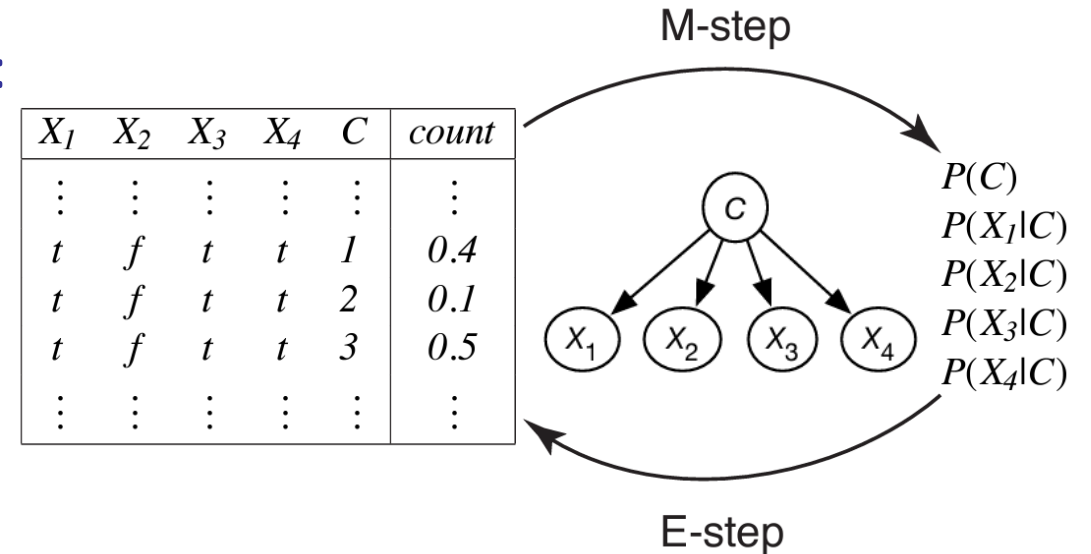
F_1	F_2	F_3	Y
$+f_1$	$+f_2$	$-f_3$	$+y$
$+f_1$	$-f_2$	$-f_3$	$+y$
$-f_1$	$+f_2$	$-f_3$	$+y$
$+f_1$	$+f_2$	$+f_3$	$-y$
$-f_1$	$+f_2$	$+f_3$	$-y$
$-f_1$	$+f_2$	$-f_3$	$-y$

$P(+f_1 +y) = \frac{2+1}{3+2} = \frac{3}{5}$	$P(+f_2 +y) = \frac{2+1}{3+2} = \frac{3}{5}$	$P(+f_3 +y) = \frac{0+1}{3+2} = \frac{1}{5}$
$P(+f_1 -y) = \frac{1+1}{3+2} = \frac{2}{5}$	$P(+f_2 -y) = \frac{3+1}{3+2} = \frac{4}{5}$	$P(+f_3 -y) = \frac{2+1}{3+2} = \frac{3}{5}$

$$\widehat{P(X)} = \frac{n(x) + \alpha}{N + \alpha|X|}$$

Expectation-Maximization

- As with HMM learning, our data may be incomplete due to *hidden* or *latent* variables
- But still important for reducing the number of parameters and data needed to learn
- Recall the *expectation-maximization* approach:
- **Expectation:** Use *inference* with current parameters to compute *expected* counts
- **Maximization:** Use expected (and observed) counts to compute new parameters
- Can initialize with some prior or guess of the inferred probabilities
- EM will converge to a local maximum of the data likelihood



Example: Expectation-Maximization

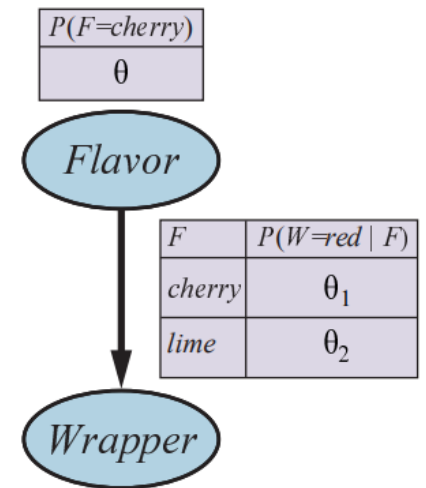
- Suppose we currently have $P(F) = (0.6, 0.4)$, $P(W = \text{red} | F) = (0.6, 0.4)$, and $P(W = \text{green} | F) = (0.4, 0.6)$

- Suppose *Flavor* is hidden in the data set, but *Wrapper* is accessible
- Inference computation of *Flavor* conditioned on *Wrapper* values:

$$P(F | \text{red}) \propto P(F)P(\text{red} | F) = \begin{pmatrix} 0.6 \\ 0.4 \end{pmatrix} \times \begin{pmatrix} 0.6 \\ 0.4 \end{pmatrix} \propto \begin{pmatrix} 0.69 \\ 0.31 \end{pmatrix}$$

$$P(F | \text{green}) \propto P(F)P(\text{green} | F) = \begin{pmatrix} 0.6 \\ 0.4 \end{pmatrix} \times \begin{pmatrix} 0.4 \\ 0.6 \end{pmatrix} \propto \begin{pmatrix} 0.5 \\ 0.5 \end{pmatrix}$$

- These are the *estimated frequencies* of the different *Flavor* values in our data set!



Example: Expectation-Maximization

- Suppose our data set contains 545 red and 455 green wrapper instances

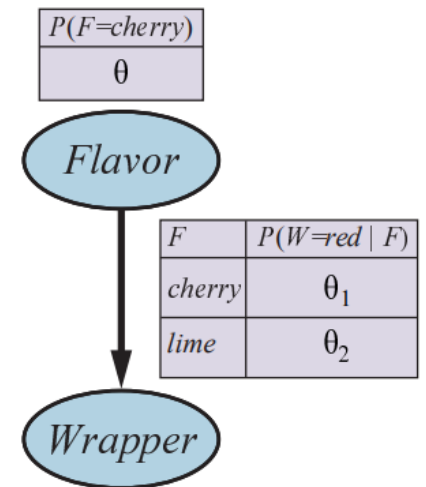
$$P(F|red) = \begin{pmatrix} 0.69 \\ 0.31 \end{pmatrix} \quad P(F|green) = \begin{pmatrix} 0.5 \\ 0.5 \end{pmatrix}$$

<i>Flavor</i>	<i>Wrapper</i>	<i>N</i>
<i>cherry</i>	<i>red</i>	376.05
<i>cherry</i>	<i>green</i>	227.5
<i>lime</i>	<i>red</i>	168.95
<i>lime</i>	<i>green</i>	227.5

- “Augmented” data set (E-step):

- New parameters (M-step): $P(F) = \frac{1}{1000} \begin{pmatrix} 376.05 + 227.5 \\ 168.95 + 227.5 \end{pmatrix} = \begin{pmatrix} 0.604 \\ 0.396 \end{pmatrix}$

$$P(W|F = \text{cherry}) = \frac{1}{603.55} \begin{pmatrix} 376.05 \\ 227.5 \end{pmatrix} = \begin{pmatrix} 0.623 \\ 0.377 \end{pmatrix} \quad P(W|F = \text{lime}) = \frac{1}{396.45} \begin{pmatrix} 168.95 \\ 227.5 \end{pmatrix} = \begin{pmatrix} 0.426 \\ 0.574 \end{pmatrix}$$



EM Considerations

- Inference can be done analytically or approximately (sampling)
- We can even *interleave* MCMC and EM iterations by treating generated samples as complete observation, giving us *approximate* expectations
- As with Baum-Welch, data likelihood is guaranteed to increase in each iteration, converging at a local maximum
- However, improvements tend to slow down after a few iterations
- Can combine with local search or gradient-based methods

Summary

- Statistical learning involves learning a hypothesis to approximate a hidden function
- Can be used for prediction, inference, pattern detection, and other tasks
- Many tradeoffs involving bias and variance, flexibility and generalizability, expressiveness and interpretability, and complexity of learning
- Regularization can help generalize models at the expense of training accuracy
- MLE is a standard technique for learning Bayes nets with complete data
- Smoothing can help regularize naïve Bayes model learning
- EM can be used when data contains hidden variables